

# Beyond Short String Optimisation: The Merits of Fixed-Length Strings

Jamie Taylor  
[jtaylor91@protonmail.com](mailto:jtaylor91@protonmail.com)

# So... `std::string`... It's pretty nice!

- Class-based representation of a string: `std::string mystring( "Hello C++ London!" );`
  - Easier and safer to work with than a C string
- This is much better than a C string! :D
  - `void foo(const char*, size_t)` becomes `void foo(const std::string&) etc.`
  - Memory allocation is handled for you
  - Short String Optimisation (SSO) buffer: short strings go on the stack

# But it's not perfect...

- SSO is **implementation-defined**
- Absence of an SSO buffer will result in **dynamic allocation**
- Strings too long for the SSO buffer will result in **dynamic allocation**
- `std::string` instance size is unavoidable, regardless of actual string length
- Non-optimal cache utilisation
- Some use-cases could benefit from an **alternative...**



# Introducing: fl::string

- A fixed-length string implementation

```
class std::string {  
    // ...  
    size_t length, capacity;  
    char sso[16]; // union with char*  
};
```

```
template<size_t length>  
class fl::string {  
    // ...  
    char m_data[length];  
};
```

length: 3								capacity: 15							
H	I	!	\0												

H	I	!	\0					4
---	---	---	----	--	--	--	--	---

- Interoperability with std::string and C strings using std::string\_view
  - inline operator string\_view() const noexcept { return string\_view{m\_data, length()}; }

# Test: Common string operations

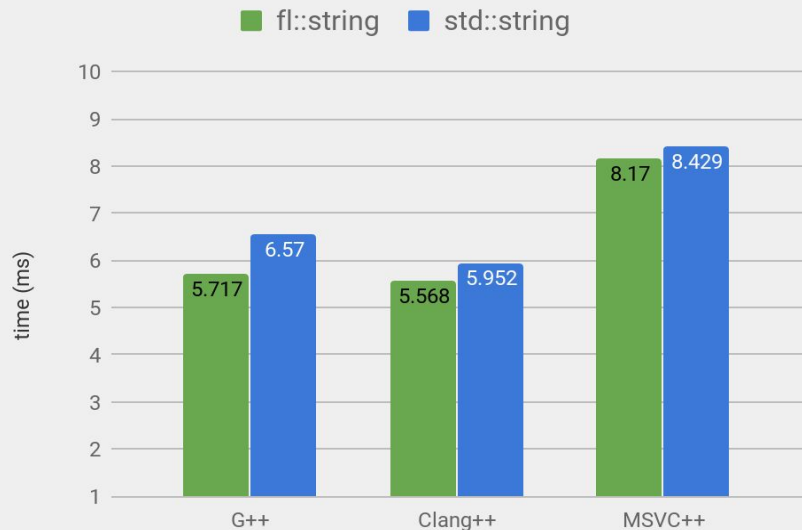
- % improvement in string operations (8 and 32 character):
  - Construction: `string::string( const char *c )` - 259% and 1049%
  - Assignment: `string::operator=( const char *c )` - 939% and 1750%
  - Concatenate: `string::operator+=( const char *c )` - 668% and 1360%
  - Avg. 6x and 13x speedup respectively

# Test: `std::unordered_map`


- Take a `std::unordered_map<>` of 128 `<string, unsigned int>` pairs
  - Key is a 4 or 8 character `std::string` or `fl::string`
  - Value represents the number of times the key has been searched for
- Record the time taken to:
  - 0) Create the map
  - 1) Perform 256 randomly-generated look-ups

# Results: std::unordered\_map - 4 char keys

std::unordered\_map - Creation (4 character keys)




std::unordered\_map - Access (4 character keys)

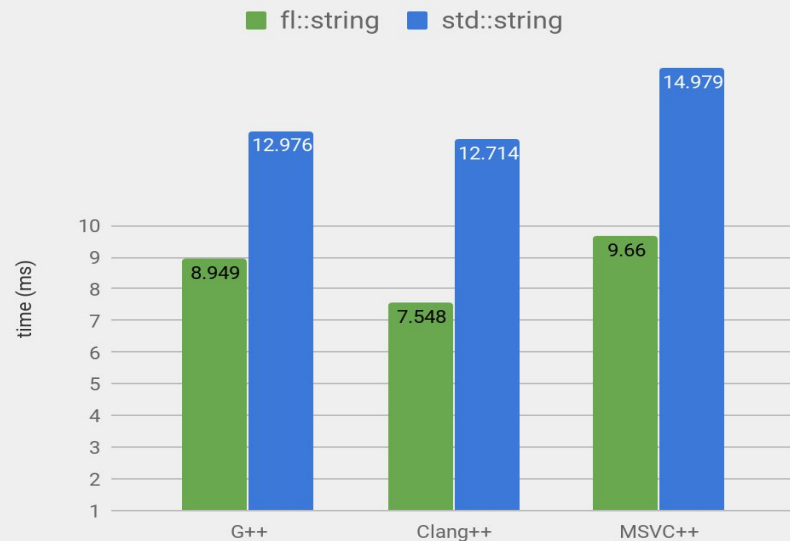


# Results: `std::unordered_map` - 8 char keys

`std::unordered_map` - Creation (8 character keys)



`std::unordered_map` - Access (8 character keys)





# Conclusions

- `std::string` isn't the be-all-end-all
  - May dynamically allocate
  - Even a 1 char string will be `sizeof(std::string)` in size
- `fl::string` offers complete control over memory layout
  - No dynamic allocation
  - Improved cache-coherency
  - Last byte free-space/null-terminator double-up
- Demonstrable **performance gains** in real-world scenarios
  - String ops: construct, assign and concatenate
  - Containers: unordered maps
- Great fit for:
  - Embedded, Games, Trading - **low-latency**

# Thank you!

[jtaylor91@protonmail.com](mailto:jtaylor91@protonmail.com)

<https://gitlab.com/jt396/flstring>



# References

- [https://david-grs.github.io/inplace\\_containers\\_for\\_fun\\_and\\_profit/](https://david-grs.github.io/inplace_containers_for_fun_and_profit/)
- <https://shaharmike.com/cpp/std-string/>
- <https://www.facebook.com/Engineering/videos/10151029396848109>

# Appendix: Testing Methodology

- % improvement in string operations (8 and 32 character):
  - Construction: `string::string( const char *c )` - 259% and 1049%
  - Assignment: `string::operator=( const char *c )` - 939% and 1750%
  - Concatenate: `string::operator+=( const char *c )` - 668% and 1360%
- Avg. 6x and 13x speedup
- Performed with 8 and 32 character strings (7 and 31 + null-terminator)
- Compilers Used: Clang++5.0.0 G++7.0.0 (MSVC data incomplete? See git.)

# Appendix: Testing Methodology cont.

- Take a `std::unordered_map<>` of 128 `<string, unsigned int>` pairs
  - Value is simply the number of times the key has been searched for - nothing special
- Perform N number of look-up operations on said maps
  - We're going to use 256
  - Look-ups are generated randomly and the same sequence is used for both the `fl::string` and `std::string` maps
- Record the time taken to:
  - 0) Create the map
  - 1) Perform the specified number of look-ups
- Repeat this M times
  - We're going to go for 1024 times
  - So: record the time taken to create the map and perform 256 look-ups 1024 times and average
  - Both 'fixed-length->std::string' and 'std::string->fixed-length' orderings used for fairness
- Performed with 4 and 8 character strings (3 and 7 + null-terminator)
- Compilers Used: Clang++5.0.0, G++7.0.0 and MSVC++2017 (19.10.25019)